

## Project Name: Drinking Water Potability Prediction using ML and H2O Auto ML



### Context :

Access to safe drinking water is essential to health, a basic human right, and a component of effective policy for health protection. This is important as a health and development issue at a national, regional, and local level. In some regions, it has been shown that investments in water supply and sanitation can yield a net economic benefit, since the reductions in adverse health effects and health care costs outweigh the costs of undertaking the interventions.

The `drinkingwaterpotability.csv` file contains water quality metrics for 3276 different water bodies

We will use different ML models and H2O Auto ML library in this project

## Time Line of the Project:

- Importing Libraries and DataSet
- Data Analysis and Preprocessing
- Feature Engineering
- Model Building using ML
- Model Building and Prediction using H2O Auto ML

### ▼ Importing Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
%matplotlib inline
```

### ▼ Loading the Data Set

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
df= pd.read_csv("/content/drive/MyDrive/drinking_water_potability.csv")
```

```
df.head()
```

|   | ph       | Hardness   | Solids      | Chloramines | Sulfate    | Conductivity | Organic_c |
|---|----------|------------|-------------|-------------|------------|--------------|-----------|
| 0 | NaN      | 204.890456 | 20791.31898 | 7.300212    | 368.516441 | 564.308654   | 10.3      |
| 1 | 3.716080 | 129.422921 | 18630.05786 | 6.635246    | NaN        | 592.885359   | 15.1      |
| 2 | 8.099124 | 224.236259 | 19909.54173 | 9.275884    | NaN        | 418.606213   | 16.8      |
| 3 | 8.316766 | 214.373394 | 22018.41744 | 8.059332    | 356.886136 | 363.266516   | 18.4      |
| 4 | 9.092223 | 181.101509 | 17978.98634 | 6.546600    | 310.135738 | 398.410813   | 11.5      |

```
df.shape
```

(3276, 10)

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3276 entries, 0 to 3275
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   ph                     2785 non-null   float64
1   Hardness               3276 non-null   float64
2   Solids                 3276 non-null   float64
3   Chloramines            3276 non-null   float64
4   Sulfate                2495 non-null   float64
5   Conductivity           3276 non-null   float64
6   Organic_carbon         3276 non-null   float64
7   Trihalomethanes        3114 non-null   float64
8   Turbidity              3276 non-null   float64
9   Potability             3276 non-null   int64   
dtypes: float64(9), int64(1)
memory usage: 256.1 KB
```

df.nunique()

```
ph                2785
Hardness          3276
Solids            3276
Chloramines       3276
Sulfate           2495
Conductivity      3276
Organic_carbon    3276
Trihalomethanes   3114
Turbidity         3276
Potability         2
dtype: int64
```

## ▼ Data Analysis

```
sns.countplot(data=df,x=df.Potability)
df.Potability.value_counts()
```

```
0    1998
1    1278
Name: Potability, dtype: int64
```



```
df.isnull().sum()
```

```
ph                491
Hardness           0
Solids             0
Chloramines        0
Sulfate           781
Conductivity       0
Organic_carbon     0
Trihalomethanes   162
Turbidity          0
Potability         0
dtype: int64
```

## ► Handling Null Values

```
[ ] ↳ 13 cells hidden
```

## ► Feature Engineering

```
[ ] ↳ 6 cells hidden
```

## ► Let us Standardize our data

```
[ ] ↳ 4 cells hidden
```

## ► Our data is ready for model building

```
[ ] ↳ 1 cell hidden
```

## ▼ Model Development

We will use the following models:

- Logistic Regression
- SVM
- Random Forest

```
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score
```

## ▼ Logistic Regression

```
lr = LogisticRegression()
lr.fit(X_train, y_train)
y_train_hat = lr.predict(X_train)
y_test_hat = lr.predict(X_test)

print('Test performance')
print('-----')
print(classification_report(y_test, y_test_hat))

print('Roc_auc score')
print('-----')
print(roc_auc_score(y_test, y_test_hat))
print('')

print('Confusion matrix')
print('-----')
print(confusion_matrix(y_test, y_test_hat))
print('')

print('accuracy score')
print('-----')
print("test data accuracy score:", accuracy_score(y_test, y_test_hat)*100)
print("train data accuracy score:", accuracy_score(y_train, y_train_hat)*100)
```

Test performance

```
-----
              precision    recall  f1-score   support

     0           0.62       1.00      0.77        610
     1           0.00       0.00      0.00        373

 accuracy               0.62            983
 macro avg           0.31            983
weighted avg           0.38            983
```

Roc\_auc score

```
-----
0.49918032786885247
```

Confusion matrix

```
-----
[[609   1]
 [373   0]]
```

accuracy score

```
-----
test data accuracy score: 61.953204476093596
train data accuracy score: 60.61927605756651
```

## ▼ Support Vector Machines

```
svm = SVC()
svm.fit(X_train, y_train)
y_train_hat = svm.predict(X_train)
y_test_hat = svm.predict(X_test)

print('Test performance')
print('-----')
print(classification_report(y_test, y_test_hat))

print('Roc_auc score')
print('-----')
print(roc_auc_score(y_test, y_test_hat))
print('')

print('Confusion matrix')
print('-----')
print(confusion_matrix(y_test, y_test_hat))
print('')

print('accuracy score')
print('-----')
print(accuracy_score(y_test, y_test_hat)*100)
print("test data accuracy score:",accuracy_score(y_test, y_test_hat)*100)
print("train data accuracy score:",accuracy_score(y_train, y_train_hat)*100)
```

Test performance

```
-----
              precision    recall  f1-score   support

     0           0.62       1.00      0.76        610
     1           0.33       0.00      0.01        373

 accuracy                   0.62        983
 macro avg              0.48       0.50      0.39        983
weighted avg              0.51       0.62      0.48        983
```

Roc\_auc score

```
-----
0.4997011383114315
```

Confusion matrix

```
-----
[[608   2]
 [372   1]]
```

accuracy score

```
-----
61.953204476093596
test data accuracy score: 61.953204476093596
train data accuracy score: 60.706498037505455
```

## ▼ Random Forest

```
rf = RandomForestClassifier(n_jobs=-1,random_state=123)
rf.fit(X_train, y_train)
y_train_hat = rf.predict(X_train)
y_test_hat = rf.predict(X_test)

print('Test performance')
print('-----')
print(classification_report(y_test, y_test_hat))

print('Roc_auc score')
print('-----')
print(roc_auc_score(y_test, y_test_hat))
print('')

print('Confusion matrix')
print('-----')
print(confusion_matrix(y_test, y_test_hat))
print('')

print('accuracy score')
print('-----')
print("test data accuracy score:",accuracy_score(y_test, y_test_hat)*100)
print("train data accuracy score:",accuracy_score(y_train, y_train_hat)*100)
```

Test performance

```
-----
              precision    recall  f1-score   support

     0           0.69       0.90       0.78        610
     1           0.66       0.33       0.44        373

 accuracy                   0.68        983
 macro avg           0.68       0.61       0.61        983
weighted avg           0.68       0.68       0.65        983
```

Roc\_auc score

```
-----
0.6140596844372171
```

Confusion matrix


```
-----
[[548  62]
 [250 123]]
```

accuracy score

```
-----
test data accuracy score: 68.26042726347914
train data accuracy score: 100.0
```

## Using Auto ML

## ▼ H2O Auto ML



H2O is a fully open-source, distributed in-memory machine learning platform with linear scalability. H2O supports the most widely used statistical & machine learning algorithms, including gradient boosted machines, generalized linear models, deep learning, and many more.

## ▼ Installing H2O Auto ML

```
!pip install requests
!pip install tabulate
!pip install "colorama>=0.3.8"
!pip install future
```



```

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/
Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (2
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-pac
Requirement already satisfied: urllib3!=1.25.0,!>=1.25.1,<1.26,>=1.21.1 in /usr/local,
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-pa
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/
Requirement already satisfied: tabulate in /usr/local/lib/python3.7/dist-packages (0
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/
Collecting colorama>=0.3.8
  Downloading colorama-0.4.5-py2.py3-none-any.whl (16 kB)
Installing collected packages: colorama
Successfully installed colorama-0.4.5
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/
Requirement already satisfied: future in /usr/local/lib/python3.7/dist-packages (0.16

```



```
!pip install h2o
```

```

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/
Collecting h2o
  Downloading h2o-3.36.1.4.tar.gz (177.1 MB)
    |████████████████████████████████████████| 177.1 MB 23 kB/s
Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (fr
Requirement already satisfied: tabulate in /usr/local/lib/python3.7/dist-packages (fr
Requirement already satisfied: future in /usr/local/lib/python3.7/dist-packages (from
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-pac
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-pa
Requirement already satisfied: urllib3!=1.25.0,!>=1.25.1,<1.26,>=1.21.1 in /usr/local,
Building wheels for collected packages: h2o
  Building wheel for h2o (setup.py) ... done
  Created wheel for h2o: filename=h2o-3.36.1.4-py2.py3-none-any.whl size=177128127 st
  Stored in directory: /root/.cache/pip/wheels/02/f9/18/5fbae4db903beda26f764b6e035ct
Successfully built h2o
Installing collected packages: h2o
Successfully installed h2o-3.36.1.4

```



## ➤ Importing the h2o Python module and H2OAutoML class

```

import h2o
from h2o.automl import H2OAutoML
h2o.init(max_mem_size='16G') ## the h2o.init() makes sure that no prior instance of H2O

```

```

Checking whether there is an H2O instance running at http://localhost:54321 ..... not
Attempting to start a local H2O server...
  Java Version: openjdk version "11.0.16" 2022-07-19; OpenJDK Runtime Environment (bu
  Starting server from /usr/local/lib/python3.7/dist-packages/h2o/backend/bin/h2o.jar
  Ice root: /tmp/tmpdouqfehz
  JVM stdout: /tmp/tmpdouqfehz/h2o_unknownUser_started_from_python.out
  JVM stderr: /tmp/tmpdouqfehz/h2o_unknownUser_started_from_python.err
  Server is running at http://127.0.0.1:54321
Connecting to H2O server at http://127.0.0.1:54321 ... successful.
H2O_cluster_uptime:      09 secs
H2O_cluster_timezone:    Etc/UTC
H2O_data_parsing_timezone: UTC
H2O_cluster_version:     3.36.1.4
H2O_cluster_version_age: 13 days
H2O_cluster_name:        H2O_from_python_unknownUser_ajkh13
H2O_cluster_total_nodes: 1
H2O_cluster_free_memory: 16 Gb

```

## ▼ Loading data

H2O cluster status: locked, healthy

```
df = h2o.import_file("/content/drive/MyDrive/drinking_water_potability.csv")
```

Parse progress:  | (c



```
df.head()
```

|  | ph      | Hardness | Solids  | Chloramines | Sulfate | Conductivity | Organic_carbon | Trihalomethanes T |
|--|---------|----------|---------|-------------|---------|--------------|----------------|-------------------|
|  | nan     | 204.89   | 20791.3 | 7.30021     | 368.516 | 564.309      | 10.3798        | 86.991            |
|  | 3.71608 | 129.423  | 18630.1 | 6.63525     | nan     | 592.885      | 15.18          | 56.3291           |
|  | 8.09912 | 224.236  | 19909.5 | 9.27588     | nan     | 418.606      | 16.8686        | 66.4201           |
|  | 8.31677 | 214.373  | 22018.4 | 8.05933     | 356.886 | 363.267      | 18.4365        | 100.342           |
|  | 9.09222 | 181.102  | 17979   | 6.5466      | 310.136 | 398.411      | 11.5583        | 31.998            |
|  | 5.58409 | 188.313  | 28748.7 | 7.54487     | 326.678 | 280.468      | 8.39973        | 54.9179           |
|  | 10.2239 | 248.072  | 28749.7 | 7.51341     | 393.663 | 283.652      | 13.7897        | 84.6036           |
|  | 8.63585 | 203.362  | 13672.1 | 4.56301     | 303.31  | 474.608      | 12.3638        | 62.7983           |
|  | nan     | 118.989  | 14285.6 | 7.80417     | 268.647 | 389.376      | 12.706         | 53.9288           |
|  | 11.1803 | 227.231  | 25484.5 | 9.0772      | 404.042 | 563.885      | 17.9278        | 71.9766           |

## ▼ H2O auto ml can do all the data preprocessing techniques

```
df_train,df_test= df.split_frame(ratios=[.8])
```

## ▼ Splitting the data

```

y = "Potability"  ## dependent variable
x = df.columns    ## Independent variable

```

```
x.remove(y)
```

## ▼ Defining the model

```
aml = H2OAutoML(max_runtime_secs=300,max_models = 10, seed = 10, verbosity="info", nfolds=
```

## ▼ Fitting the model

```
aml.train(x=x,y=y, training_frame=df_train)
```

AutoML progress: |  
20:46:34.983: Project: AutoML\_1\_20220816\_204634  
20:46:34.994: Setting stopping tolerance adaptively based on the training frame: 0.01  
20:46:34.994: Build control seed: 10  
20:46:34.995: training frame: Frame key: AutoML\_1\_20220816\_204634\_training\_py\_3\_sid\_3  
20:46:34.996: validation frame: NULL  
20:46:35.5: leaderboard frame: NULL  
20:46:35.6: blending frame: NULL  
20:46:35.6: response column: Potability  
20:46:35.6: fold column: null  
20:46:35.6: weights column: null  
20:46:35.113: Loading execution steps: [{XGBoost : [def\_2 (1g, 10w), def\_1 (2g, 10w)],  
20:46:35.215: AutoML job created: 2022.08.16 20:46:34.844  
20:46:35.220: AutoML build started: 2022.08.16 20:46:35.218  
20:46:35.315: AutoML: starting XGBoost\_1\_AutoML\_1\_20220816\_204634 model training  
20:46:35.333: \_response param, We have detected that your response column has only 2  
█  
20:46:40.44: New leader: XGBoost\_1\_AutoML\_1\_20220816\_204634, rmse: 0.5149273528821546  
20:46:40.96: AutoML: starting GLM\_1\_AutoML\_1\_20220816\_204634 model training  
20:46:40.107: \_response param, We have detected that your response column has only 2  
█  
20:46:41.890: New leader: GLM\_1\_AutoML\_1\_20220816\_204634, rmse: 0.4882227981301222  
20:46:41.902: AutoML: starting GBM\_1\_AutoML\_1\_20220816\_204634 model training  
20:46:41.907: \_response param, We have detected that your response column has only 2  
█  
20:46:45.315: New leader: GBM\_1\_AutoML\_1\_20220816\_204634, rmse: 0.47772445901616317  
20:46:45.317: AutoML: starting XGBoost\_2\_AutoML\_1\_20220816\_204634 model training  
20:46:45.321: \_response param, We have detected that your response column has only 2  
█  
20:46:47.231: AutoML: starting DRF\_1\_AutoML\_1\_20220816\_204634 model training  
20:46:47.232: \_response param, We have detected that your response column has only 2  
█  
20:46:54.621: New leader: DRF\_1\_AutoML\_1\_20220816\_204634, rmse: 0.47239112244862225  
20:46:54.623: AutoML: starting GBM\_2\_AutoML\_1\_20220816\_204634 model training  
20:46:54.624: \_response param, We have detected that your response column has only 2  
█  
20:46:56.178: New leader: GBM\_2\_AutoML\_1\_20220816\_204634, rmse: 0.4719602021948012  
20:46:56.180: AutoML: starting GBM\_3\_AutoML\_1\_20220816\_204634 model training  
20:46:56.180: \_response param, We have detected that your response column has only 2  
20:46:57.636: AutoML: starting GBM\_4\_AutoML\_1\_20220816\_204634 model training  
20:46:57.637: \_response param, We have detected that your response column has only 2  
█  
20:46:58.954: AutoML: starting XGBoost\_3\_AutoML\_1\_20220816\_204634 model training  
20:46:58.955: \_response param, We have detected that your response column has only 2  
20:46:59.733: AutoML: starting XRT\_1\_AutoML\_1\_20220816\_204634 model training  
20:46:59.733: \_response param, We have detected that your response column has only 2  
█  
20:47:03.397: New leader: XRT\_1\_AutoML\_1\_20220816\_204634, rmse: 0.47161026155960806

```
20:47:03.399: No base models, due to timeouts or the exclude_algos option. Skipping
20:47:03.421: AutoML: starting StackedEnsemble_BestOfFamily_1_AutoML_1_20220816_20463
20:47:03.422: _response param, We have detected that your response column has only 2
```

```
20:47:04.208: New leader: StackedEnsemble_BestOfFamily_1_AutoML_1_20220816_204634, rm
20:47:04.216: AutoML: starting StackedEnsemble_AllModels_1_AutoML_1_20220816_204634 n
20:47:04.218: _response param, We have detected that your response column has only 2
```

```
(done) 100%
```

```
20:47:05.167: Actual modeling steps: [{XGBoost : [def_2 (1g, 10w)]}, {GLM : [def_1 (1
20:47:05.168: AutoML build stopped: 2022.08.16 20:47:05.167
20:47:05.168: AutoML build done: built 10 models
20:47:05.168: AutoML duration: 29.949 sec
```

#### Model Details

```
=====
```

```
H2OStackedEnsembleEstimator : Stacked Ensemble
```

```
Model Key: StackedEnsemble_BestOfFamily_1_AutoML_1_20220816_204634
```

```
No model summary for this model
```

```
ModelMetricsRegressionGLM: stackedensemble
```

```
** Reported on train data. **
```

```
MSE: 0.07662728232735182
```

```
RMSE: 0.27681633320191174
```

## ▼ Seeing the Leaderboard

```
Mean Residual Deviance: 0.07662728232735182
```

```
lb = aml.leaderboard
```

```
Null deviance: 621.3210445468519
```

```
lb
```

| model_id  | rmse     | mse      | mae      | rr    |
|---|----------|----------|----------|-------|
| StackedEnsemble_BestOfFamily_1_AutoML_1_20220816_204634 | 0.468114 | 0.219131 | 0.439998 | 0.329 |
| StackedEnsemble_AllModels_1_AutoML_1_20220816_204634    | 0.468693 | 0.219673 | 0.43812  | 0.329 |
| XRT_1_AutoML_1_20220816_204634                          | 0.47161  | 0.222416 | 0.436527 | 0.332 |
| GBM_2_AutoML_1_20220816_204634                          | 0.47196  | 0.222746 | 0.44369  | 0.331 |
| DRF_1_AutoML_1_20220816_204634                          | 0.472391 | 0.223153 | 0.438843 | 0.332 |
| GBM_4_AutoML_1_20220816_204634                          | 0.473481 | 0.224185 | 0.440279 | 0.333 |
| GBM_3_AutoML_1_20220816_204634                          | 0.475852 | 0.226435 | 0.445579 | 0.335 |
| GBM_1_AutoML_1_20220816_204634                          | 0.477724 | 0.228221 | 0.457452 | 0.335 |
| GLM_1_AutoML_1_20220816_204634                          | 0.488223 | 0.238362 | 0.476737 | 0.343 |
| XGBoost 3 AutoML 1 20220816 204634                      | 0.491749 | 0.241817 | 0.432158 | 0.346 |

## ▼ Getting all the model ids

```
model_ids = list(aml.leaderboard['model_id'].as_data_frame().iloc[:,0])
```

```
model_ids
```

```
['StackedEnsemble_BestOfFamily_1_AutoML_1_20220816_204634',  
 'StackedEnsemble_AllModels_1_AutoML_1_20220816_204634',  
 'XRT_1_AutoML_1_20220816_204634',  
 'GBM_2_AutoML_1_20220816_204634',  
 'DRF_1_AutoML_1_20220816_204634',  
 'GBM_4_AutoML_1_20220816_204634',  
 'GBM_3_AutoML_1_20220816_204634',  
 'GBM_1_AutoML_1_20220816_204634',  
 'GLM_1_AutoML_1_20220816_204634',  
 'XGBoost_3_AutoML_1_20220816_204634',  
 'XGBoost_1_AutoML_1_20220816_204634',  
 'XGBoost_2_AutoML_1_20220816_204634']
```

```
aml.leader.model_performance(df_test)
```

```
ModelMetricsRegressionGLM: stackedensemble  
** Reported on test data. **
```

```
MSE: 0.21298542655770972  
RMSE: 0.4615034415448163  
MAE: 0.4316828902199786  
RMSLE: 0.3264671388593695  
R^2: 0.09410380451873568  
Mean Residual Deviance: 0.21298542655770972  
Null degrees of freedom: 671  
Residual degrees of freedom: 667  
Null deviance: 158.15063673752553  
Residual deviance: 143.12620664678093  
AIC: 879.7841967507071
```

## ▼ Getting the model details for best performing model

```
h2o.get_model([mid for mid in model_ids if "StackedEnsemble" in mid][0])
```

## Model Details

=====

H2OStackedEnsembleEstimator : Stacked Ensemble

Model Key: StackedEnsemble\_BestOfFamily\_1\_AutoML\_1\_20220816\_204634

No model summary for this model

ModelMetricsRegressionGLM: stackedensemble

\*\* Reported on train data. \*\*

MSE: 0.07662728232735182

RMSE: 0.27681633320191174

MAE: 0.2571869884087675

RMSLE: 0.1955125016838729

R^2: 0.6788496946438487

Mean Residual Deviance: 0.07662728232735182

Null degrees of freedom: 2603

Residual degrees of freedom: 2599

Null deviance: 621.3210445468519

Residual deviance: 199.53744318042413

AIC: 712.6712138468265

ModelMetricsRegressionGLM: stackedensemble

\*\* Reported on cross-validation data. \*\*

MSE: 0.21913074100035862

RMSE: 0.4681140256394361

MAE: 0.4399982381728644

RMSLE: 0.32933724813288806

R^2: 0.08160772184836851

Mean Residual Deviance: 0.21913074100035862

Null degrees of freedom: 2603

Residual degrees of freedom: 2599

Null deviance: 621.6678502235175

Residual deviance: 570.6164495649339

AIC: 3448.7340192803363

```
output= h2o.get_model([mid for mid in model_ids if "StackedEnsemble" in mid][0])
output.params
```

```
{'name': 'GLM_1_AutoML_1_20220816_204634',
 'type': 'Key<Keyed>',
 'URL': None},
{'__meta': {'schema_version': 3,
 'schema_name': 'KeyV3',
 'schema_type': 'Key<Keyed>'},
 'name': 'XGBoost_3_AutoML_1_20220816_204634',
 'type': 'Key<Keyed>',
 'URL': None}],
'input': [{'__meta': {'schema_version': 3,
 'schema_name': 'KeyV3',
 'schema_type': 'Key<Keyed>'},
 'name': 'XRT_1_AutoML_1_20220816_204634',
 'type': 'Key<Keyed>',
 'URL': None},
{'__meta': {'schema_version': 3,
 'schema_name': 'KeyV3',
 'schema_type': 'Key<Keyed>'},
 'name': 'GBM_2_AutoML_1_20220816_204634',
 'type': 'Key<Keyed>',
 'URL': None}]
```





\_\_\_\_\_

Model Key: StackedEnsemble BestOfFamily 1 AutoML 1 20220816 204634

```
ModelMetricsRegressionGLM: stackedensemble
```

```
MSE: 0.07662728232735182
RMSE: 0.27681633320191174
MAE: 0.2571869884087675
RMSLE: 0.1955125016838729
R^2: 0.6788496946438487
Mean Residual Deviance: 0.07662728232735182
Null degrees of freedom: 2603
Residual degrees of freedom: 2599
Null deviance: 621.3210445468519
Residual deviance: 199.53744318042413
AIC: 712.6712138468265
```

**\*\* Reported on cross-validation data. \*\***

```
MSE: 0.21913074100035862
RMSE: 0.4681140256394361
MAE: 0.4399982381728644
RMSLE: 0.32933724813288806
R^2: 0.08160772184836851
Mean Residual Deviance: 0.21913074100035862
Null degrees of freedom: 2603
Residual degrees of freedom: 2599
Null deviance: 621.6678502235175
Residual deviance: 570.6164495649339
AIC: 3448.7340192803363
```

```
stackensemble prediction progress: ██████████ (c
```



|   |                        |          |          |          |          |
|---|------------------------|----------|----------|----------|----------|
| 1 | mean residual deviance | 0.219096 | 0.004766 | 0.215726 | 0.222466 |
|---|------------------------|----------|----------|----------|----------|

y\_pred

**predict**

0.485927

0.398239

0.386352

0.350281

0.49567

0.386966

0.182781

0.274113

0.30379

0.393814